

Find a car Park

Authors: Noa Aizer & Ido Shapira & Shay Leyzerovich

In this project we propose to develop a model for identifying whether there is free space in reparking lot image, or whether the parking is full.

Report:

In this project we have classification problem. We wish to recognize whether the parking spot is full or free in an image. Our data set contains a huge amount of parking spots images. Therefore our features will be the images pixels and the label divided to free and full.

Preprocessing: Preparing Data:

In order to use the data we needed to make some adjustments.

- In the first stage we imported the data, we tried to work with RGB data but it was too heavy to handle on our computer, so we used grayscale images but we've got lower results so we decided to stay with our original images. Another thing is that the data was too big and cause hardware and memory problems- there was too many images and all of them with high quality. Thus we resized the pictures and took some samples from the data instead of all the images in the dataset.
- In the next stage we did label encoding: 1 for free, 0 for full.
- We normalized all the data in order to clean noise from the images.
- In addition, we defined a random seed in order to compare between the logistic model and the MLP model.
- After that, we divide the set into two parts: 80% for training and 20% for testing. We decided not to make validation data because of our low hardware.

Building the logistic regression model:

At first, we build a basic model for our problem and check the results.

- We improved the model by changing the function to sigmoid instead of the function we saw in class.
- In order to Faced the problem of overfitting/underfitting we reduced the number of iteration and change the learning rate until we have got a nice decreasing of our loss function.
- The logistic model got 86% correct on the testing set after training.
- According to the confusion matrix we can conclude it was more difficult for the logistic model to identify images of free parking than full parking images.

Building the MLP (Multi-Layer Perceptron) model:

- We build the network with 2 hidden layers ; the first hidden layer has 128 neurons and the second has 256 neurons.
- We used the relu activation function to bring more power to the model.
- We tuned the dropout function to avoid over fitting but the train error was very high and it extended the time of the model training, so in the end we decided not to use it.
- The MLP model got 90% correct on the testing set after training.
- In the beginning we use a low learning rate and because of that we never got to the minimum loss. In order to fix that we choose a bigger learning rate.
- We weren't satisfied with the results and in order to improve them:
 - we changed the amount of iterations - to avoid over fitting
 - we changed the size and amount of hidden layers.

Comparing to the logistic regression model:

We've saw that the MLP got better results but it is more complex model than the logistic regression, therefore it took more time to train it.

You can find more specific details, screenshots and plots during the code below.

Code & Explanations:

Import all the necessary libraries:

We have used Python libraries:

- tensorflow v1
- sklearn

```
In [1]: import tensorflow.compat.v1 as tf
import tensorflow as tf
import cv2
import os
from tqdm import tqdm
import numpy as np
import random as rs

from sklearn.model_selection import train_test_split, cross_val_score
from sklearn.linear_model import LogisticRegression
from sklearn.preprocessing import LabelEncoder
from sklearn.metrics import confusion_matrix
from sklearn import metrics

import matplotlib.pyplot as plt
import pandas as pd
import seaborn as sns

rs = 42 # random seed

fix random seed:

In [2]: np.random.seed(rs)
rs.seed(rs)
tf.set_random_seed(rs)

import the data:

In [3]: X = []
Z = []
IMG_SIZE=150

FREE_DIR='./home/ido/Desktop/Find-a-Car-Park/data/Free'
FULL_DIR='./home/ido/Desktop/Find-a-Car-Park/data/Full'

In [4]: def make_train_data(labels_DIR):
    for img in tqdm(os.listdir(DIR)):
        path = os.path.join(DIR,img)
        img = plt.imread(path)
        img = cv2.resize(img, (IMG_SIZE, IMG_SIZE))
        img = cv2.cvtColor(img, cv2.COLOR_BGR2GRAY)
        np_img=np.asarray(img)
        X.append(np_img)
        Z.append(str(label))

In [5]: # make 'Full' data
make_train_data('Full',FULL_DIR)
print(len(X))

100%|#####| 2195/2195 [00:28<00:00, 104.63it/s]
2195
False

In [6]: # make 'Free' data
make_train_data('Free',FREE_DIR)
print(len(X))

100%|#####| 1067/1067 [00:09<00:00, 109.39it/s]
1067
True
```

The description of the data:

- Number of instances: ---
- Number of attributes: --- (type here what is the feature)

```
In [7]: # check some image
fig,ax=plt.subplots(3,2)
fig.set_size_inches(15,15)
for i in range(3):
    for j in range(2):
        i=mn.randint(0,len(Z))
        ax[i,j].imshow(X[i])
        ax[i,j].set_title('Label: '+Z[i])
plt.tight_layout()
```



Label encoding of the target:

Full= 0
Free= 1

```
In [8]: Y = np.array(Z)
Y = np.where(Y=='Full', 0, Y)
Y = np.where(Y=='Free', 1, Y)
Y = Y.astype('int32')
```

normalization and flatten the images:

```
In [9]: X = np.array([cv2.normalize(X[i], None, 0, 255, cv2.NORM_MINMAX) for i in range(len(X))]) # normalization
X=np.array(X)/255
X = X.flatten().reshape(len(X), (IMG_SIZE ** 2) * 3) # 3 is because it is RGB image
# X = X.flatten().reshape(len(X), (IMG_SIZE ** 2))
```

```
In [10]: df = pd.DataFrame(data=X)
df

Out[10]:
```

	0	1	2	3	4	5	6	7	8	9	...	67490	67491	67492	67493	67494	67495	67496	67497	67498	67499
0	0.294118	0.176471	0.243137	0.278431	0.164706	0.227451	0.286275	0.176471	0.215686	0.284118	...	0.960784	0.945098	0.972549	0.223529	0.215686	0.282353	0.164706	0.145098	0.215686	
1	0.474510	0.372756	0.272549	0.490109	0.356863	0.388235	0.196161	0.333333	0.384314	0.454902	...	0.211765	0.411765	0.274510	0.258294	0.356863	0.188235	0.333333	0.188235	0.184314	
2	0.298039	0.384314	0.243137	0.313725	0.227451	0.254902	0.266667	0.160784	0.207843	0.262745	...	0.349020	0.376471	0.329412	0.231373	0.211765	0.164706	0.078431	0.376471	0.313725	0.266667
3	0.588235	0.384314	0.482353	0.670588	0.439216	0.549020	0.741176	0.484118	0.631373	0.639216	...	0.274510	0.243137	0.168827	0.219098	0.176471	0.105982	0.160784	0.270588	0.223529	
4	0.219608	0.176471	0.227451	0.235294	0.188235	0.256980	0.203922	0.164706	0.219608	0.247059	...	0.172549	0.376471	0.439216	0.435294	0.592157	0.103992	0.780392	0.792157	0.964706	0.960784
...
3257	0.517647	0.364706	0.466667	0.533333	0.368627	0.470588	0.541176	0.384314	0.482353	0.498039	...	0.364706	0.360784	0.298039	0.349020	0.360784	0.286275	0.345098	0.360784	0.274510	0.317647
3258	0.290196	0.188235	0.223529	0.282353	0.176471	0.203922	0.282353	0.203922	0.211765	0.301961	...	0.160784	0.211765	0.211765	0.282745	0.227451	0.282353	0.266667	0.254902	0.325490	
3259	0.541176	0.407843	0.458824	0.545098	0.403922	0.454902	0.521569	0.372549	0.431373	0.498039	...	0.325490	0.301961	0.227451	0.278431	0.258824	0.176471	0.247059	0.207843	0.113725	0.164706
3260	0.392157	0.278431	0.356863	0.470588	0.333333	0.407843	0.443137	0.321569	0.458824	...	0.321569	0.352941	0.298039	0.317647	0.303161	0.258824	0.206667	0.258824	0.211765	0.215686	
3261	0.101961	0.227451	0.074510	0.128412	0.054902	0.105882	0.172549	0.109804	0.152941	0.192157	...	0.349020	0.364706	0.333333	0.345098	0.356863	0.317647	0.341176	0.341176	0.298039	0.309804

3262 rows * 67500 columns

We can see that the size of an image is 150 * 150 * 3 = 67500 that later it would be our features in our models.

```
In [11]: X.shape
Out[11]: (3262, 67500)
```

Separate Split:

```
In [12]: # separate data
X_train,X_test,y_train,y_test=train_test_split(X,Y,test_size=0.20,random_state=rs, shuffle=True)
freqs = pd.DataFrame({"Training dataset": [(y_train == 1).sum(), (y_train == 0).sum()],
                    "Test dataset": [(y_test == 1).sum(), (y_test == 0).sum()],
                    "Total": [(y_train == 1).sum() + (y_test == 1).sum(), (y_train == 0).sum() + (y_test == 0).sum()]},
                    index=["Free", "Full"])
freqs[["Training dataset", "Test dataset", "Total"]]

Out[12]:
```

	Training dataset	Test dataset	Total
Free	863	204	1067
Full	1746	449	2195

We'll build a classifier with two classes: "full", so we create the according labels.

```
In [13]: labels_train = (np.arange(2) == y_train[:,None]).astype(np.float32)
labels_test = (np.arange(2) == y_test[:,None]).astype(np.float32)
```

Define a function that calculate the statistics:

```
In [14]: def test_statistics(confusion):
    TP = sess.run(confusion[1, 1])
    TN = sess.run(confusion[0, 0])
    FP = sess.run(confusion[0, 1])
    FN = sess.run(confusion[1, 0])
    print('True Positives (TP):', TP)
    print('True Negatives (TN):', TN)
    print('False Positives (FP):', FP)
    print('False Negatives (FN):', FN)
    print('Classification Accuracy:', (TP + TN) / float(TP + TN + FP + FN))
    precision = TP / float(TP + FP)
    recall = TP / float(TP + FN)
    avg_loss = batch_loss / recall
    print('f-score:', 2 * precision * recall / (precision + recall))

1. Declare the model you need to learn
```

```
In [15]: tf.compat.v1.disable_eager_execution()
data = tf.placeholder(dtype=tf.float32, shape=[None, X.shape[1]])
target = tf.placeholder(dtype=tf.float32, shape=[None, 1])

1. Declare the variables that need to be learned and initialization.
```

The number of variables is the number of the features (X.shape[1] 1)

```
In [16]: W = tf.Variable(tf.zeros(shape=[X.shape[1], 1])) # tf.random_normal maybe better
b1 = tf.Variable(tf.zeros(shape=[1, 1]))
init = tf.global_variables_initializer()
sess = tf.Session()
sess.run(init)

1. Declare the model you need to learn
```

```
In [17]: mod = tf.add(tf.matmul(data, W), b)
# y = 1 / (1.0 + tf.exp(-tf.matmul(data, W) + b)) # tf.nn.sigmoid(tf.matmul(data, W) + b)

1. Define the requires parameters:
```

- learning rate
- batch size
- number of iteration

```
In [18]: learning_rate = 0.003
batch_size = 128
iter_num = 850

1. Loss function and optimizer:
```

We use the sigmoid and cross-entropy loss function,

meaning, first doing a sigmoid on the model result and then using the cross-entropy loss function.

```
In [19]: # We want to minimize the loss function using the Gradient-Descent method
loss = tf.reduce_mean(tf.nn.sigmoid_cross_entropy_with_logits(logits=mod, labels=target))
optimizer = tf.train.GradientDescentOptimizer(learning_rate).minimize(loss)

1. Define the accuracy
```

```
In [20]: # The default threshold is 0.5, rounded off directly
prediction = tf.round(tf.sigmoid(mod))
correct = tf.cast(tf.equal(prediction, target), dtype=tf.float32)
# Average
accuracy = tf.reduce_mean(correct)
```

End of the definition of the model framework

```
In [21]: # Define the variable that saves the result
loss_trace = []
train_acc = []
test_acc = []

Start training the Logistic-Regression model:
```

```
In [22]: for epoch in range(iter_num):
    # Generate random batch index
    batch_index = np.random.choice(len(X_train), size=batch_size)
    batch_train_X = X_train[batch_index]
    batch_train_y = np.matrix(y_train[batch_index]).T
    sess.run(optimizer, feed_dict={data: batch_train_X, target: batch_train_y})
    temp_loss = sess.run(loss, feed_dict={data: batch_train_X, target: batch_train_y})

    # convert into a matrix, and the shape of the placeholder to correspond
    temp_train_acc = sess.run(accuracy, feed_dict={X_train, target: np.matrix(y_train).T})
    temp_test_acc = sess.run(accuracy, feed_dict={data: X_test, target: np.matrix(y_test).T})

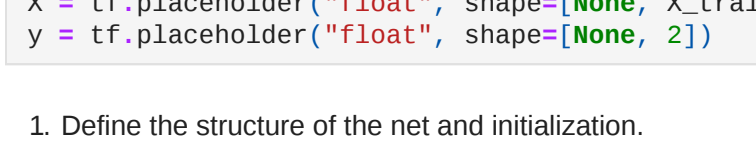
    # recode the result
    loss_trace.append(temp_loss)
    train_acc.append(temp_train_acc)
    test_acc.append(temp_test_acc)

    # output
    if (epoch + 1) % 50 == 0:
        print('epoch: {:4d} loss: {:.5f} train_acc: {:.5f} test_acc: {:.5f}'.format(epoch + 1, temp_loss,
                                          temp_train_acc, temp_test_acc))

epoch: 50 loss: 0.435215 train_acc: 0.773869 test_acc: 0.779479
epoch: 100 loss: 0.468824 train_acc: 0.792841 test_acc: 0.797858
epoch: 150 loss: 0.359847 train_acc: 0.817171 test_acc: 0.803982
epoch: 200 loss: 0.408420 train_acc: 0.752779 test_acc: 0.773354
epoch: 250 loss: 0.344399 train_acc: 0.842852 test_acc: 0.828484
epoch: 300 loss: 0.332117 train_acc: 0.842852 test_acc: 0.854518
epoch: 350 loss: 0.279847 train_acc: 0.878251 test_acc: 0.869643
epoch: 400 loss: 0.334233 train_acc: 0.884639 test_acc: 0.857580
epoch: 450 loss: 0.309067 train_acc: 0.894979 test_acc: 0.868308
epoch: 500 loss: 0.231778 train_acc: 0.877748 test_acc: 0.874426
epoch: 550 loss: 0.254588 train_acc: 0.845918 test_acc: 0.846861
epoch: 600 loss: 0.268418 train_acc: 0.888031 test_acc: 0.889551
epoch: 650 loss: 0.236767 train_acc: 0.903795 test_acc: 0.894334
epoch: 700 loss: 0.285662 train_acc: 0.915293 test_acc: 0.901991
epoch: 750 loss: 0.241730 train_acc: 0.924875 test_acc: 0.893865
epoch: 800 loss: 0.240127 train_acc: 0.927942 test_acc: 0.889928
epoch: 850 loss: 0.236828 train_acc: 0.927688 test_acc: 0.894334

Visualization of the loss function results:
```

```
In [23]: plt.plot(loss_trace)
plt.title('Cross Entropy Loss')
plt.xlabel('epoch')
plt.ylabel('loss')
plt.show()
```



We can see that the loss value reduce during the train, there's a graph that shows that.

```
In [24]: y_predictions = sess.run(prediction, feed_dict={data: X_test, target: np.matrix(y_test).T})

In [25]: tf.compat.v1.disable_eager_execution()
confusion = tf.math.confusion_matrix(y_test, y_predictions, num_classes=2, dtype=tf.int32)
print(sess.run(confusion))

[[410 39]
 [ 30 174]]
```

```
In [26]: statistics(confusion)
True Positives (TP): 174
True Negatives (TN): 410
False Positives (FP): 39
False Negatives (FN): 30
Classification Accuracy: 0.8943338437978561
Precision: 0.81549118048889
Recall: 0.8529411764705882
f-score: 0.882817636827215
```

The results indicate that the model could not recognize a lot of free parking spot because of the recall but when it did the precision was high.

Begin building the MLP model using tensorflow.v1:

```
1. Define placeholders

In [68]: tf.compat.v1.disable_eager_execution()
X = tf.placeholder(dtype=tf.float, shape=[None, X_train.shape[1]])
y = tf.placeholder(dtype=tf.float, shape=[None, 2])

1. Define the structure of the net and initialization.
```

```
In [61]: # Layer's sizes
x_size = X_train.shape[1] # The number of pixels in the image
y_size = 2 # Number of outcomes ('Free' or 'Full')

# First layer
hid1_size = 256
w1 = tf.Variable(tf.random_normal([hid1_size, x_size], stddev=0.01), name='w1')
b1 = tf.Variable(tf.constant(0.1, shape=[hid1_size, 1]), name='b1')
z1 = tf.nn.dropout(tf.nn.relu(tf.matmul(w1, tf.transpose(X)), b1), keep_prob=0.5)
y1 = tf.nn.relu(tf.add(tf.matmul(w1, tf.transpose(X)), b1))

# Second layer
hid2_size = 128
w2 = tf.Variable(tf.random_normal([hid2_size, hid1_size], stddev=0.01), name='w2')
b2 = tf.Variable(tf.constant(0.1, shape=[hid2_size, 1]), name='b2')
z2 = tf.nn.dropout(tf.nn.relu(tf.add(tf.matmul(w2, y1), b2)), keep_prob=0.5)
y2 = tf.nn.relu(tf.add(tf.matmul(w2, y1), b2))

# Output layer
wo = tf.Variable(tf.random_normal([y_size, hid2_size], stddev=0.01), name='wo')
bo = tf.Variable(tf.random_normal([y_size, 1]), name='bo')
yo = tf.transpose(tf.add(tf.matmul(wo, y2), bo))

sess = tf.Session()
init = tf.global_variables_initializer()
sess.run(init)
```

1. Define the requires parameters:

- learning rate
- batch size
- number of iteration

```
In [62]: learning_rate = 0.003
batch_size = 128
iter_num = 650

1. Loss function and optimizer (Back Propagation):
```

```
In [63]: loss = tf.reduce_mean(tf.nn.softmax_cross_entropy_with_logits(logits=yo))
updates = tf.train.GradientDescentOptimizer(learning_rate).minimize(loss)

1. Define the accuracy
```

```
In [64]: predit = tf.argmax(yo, axis=1)

End of the definition of the model framework
```

```
In [65]: # Define the variable that saves the result
loss_trace = []
train_acc = []
test_acc = []

Start training the MLP model:
```

```
In [66]: for epoch in range(iter_num):
    avg_loss = 0.0

    # Generate random batch index
    batch_index = np.random.choice(len(X_train), size=batch_size)
    batch_train_X = X_train[batch_index]
    batch_train_y = np.matrix(labels_train[batch_index])

    for i in range(batch_train_X.shape[0]):
        sess.run(updates, feed_dict={X: batch_train_X[i + 1], y: batch_train_y[i + 1]})
        temp_loss = sess.run(loss, feed_dict={X: batch_train_X[i, None], y: batch_train_y[i, None]})
        avg_loss += temp_loss

    # convert into a matrix, and the shape of the placeholder to correspond
    temp_train_acc = np.mean(np.argmax(labels_train, axis=1) == sess.run(predit, feed_dict={X: X_train, y: labels_test}))
    temp_test_acc = np.mean(np.argmax(labels_test, axis=1) == sess.run(predit, feed_dict={X: X_test, y: labels_test}))
    avg_loss /= batch_train_X.shape[0]


    # recode the result
    loss_trace.append(avg_loss)
    train_acc.append(temp_train_acc)
    test_acc.append(temp_test_acc)

    # output
    if (epoch + 1) % 50 == 0:
        print('epoch: {:4d} loss: {:.5f} train_acc: {:.5f} test_acc: {:.5f}'.format(epoch + 1, avg_loss, temp_train_acc, temp_test_acc))

epoch: 50 loss: 0.092926 train_acc: 0.874215 test_acc: 0.857580
epoch: 100 loss: 0.064709 train_acc: 0.954909 test_acc: 0.915773
epoch: 150 loss: 0.048524 train_acc: 0.972829 test_acc: 0.941897
epoch: 200 loss: 0.014883 train_acc: 0.972787 test_acc: 0.955590
epoch: 250 loss: 0.01425 train_acc: 0.972829 test_acc: 0.959121
epoch: 300 loss: 0.006016 train_acc: 0.981682 test_acc: 0.958652
epoch: 350 loss: 0.004346 train_acc: 0.987735 test_acc: 0.960184
epoch: 400 loss: 0.001360 train_acc: 0.993867 test_acc: 0.964778
epoch: 450 loss: 0.003090 train_acc: 0.993867 test_acc: 0.967841
epoch: 500 loss: 0.003791 train_acc: 0.935991 test_acc: 0.988116
epoch: 550 loss: 0.002040 train_acc: 0.990934 test_acc: 0.967841
epoch: 600 loss: 0.003392 train_acc: 0.988591 test_acc: 0.960184

Visualization of the loss function results:
```

```
In [67]: plt.plot(loss_trace)
plt.title('Cross Entropy Loss')
plt.xlabel('epoch')
plt.ylabel('loss')
plt.show()
```



```
In [68]: y_predictions = sess.run(predit, feed_dict={X: X_test, y: labels_test})

In [69]: confusion = tf.math.confusion_matrix(y_test, y_predictions, num_classes=2, dtype=tf.int32)
print(sess.run(confusion))

[[439 19]
 [ 7 197]]
```

```
In [70]: statistics(confusion)
True Positives (TP): 197
True Negatives (TN): 439
False Positives (FP): 19
False Negatives (FN): 7
Classification Accuracy: 0.9601837672281777
Precision: 0.9576937416481069
Recall: 0.965686745898939
f-score: 0.96168369877978
```